

- **Uninformed search** 无信息的搜索：除了问题中提供的定义之外没有任何关于状态的附加信息。
- **Informed search** 有信息的搜索：在问题本身的定义之外还可利用问题的特定知识。

# Review: Last Class

2

```
function TREE-SEARCH( problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INTIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT (fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds
      return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

A strategy is defined by picking the **order of node expansion**

Variety of **uninformed** search strategies

Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

# Uninformed Search Strategies

3

**Uninformed** search strategies use only the information available in the problem definition

- Breadth-first search (广度优先搜索)
- Uniform-cost search (代价一致搜索)
- Depth-first search (深度优先搜索)
- Depth-limited search (深度有限搜索)
- Iterative deepening search (迭代深入深度优先搜索)

# Uninformed Search Strategies

4

- Summary of uninformed tree (problem-solving) search:
  - ▣ Algorithms differ by method of *expansion*, or choice of *state-action* sequence for offline evaluation
  - ▣ Complexity tradeoffs, but poor (worst case or typical case) performance from all when state space is large

# 5 Informed Search Algorithms

## Chapter 4



# Outline

6

- **Best-first search** (最佳优先搜索)
  - Greedy best-first search
  - A\* search
- Heuristics
- Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms

# Best-First Search

7

**Idea:** use an **evaluation function** (评价函数)  $f(n)$  for each node

— estimate of “desirability”

⇒ Expand most desirable unexpanded node

**Implementation:**

**fringe** is a queue sorted in decreasing order of desirability

— priority queue (优先级队列)

**Special cases:**



greedy search

A\* search

# Uniform Cost

8

- Strategy: expand lowest path cost
- The good: UCS is complete and optimal!
- The bad:
  - ▣ Explores options in every “direction”
  - ▣ No information about goal location







# Best First

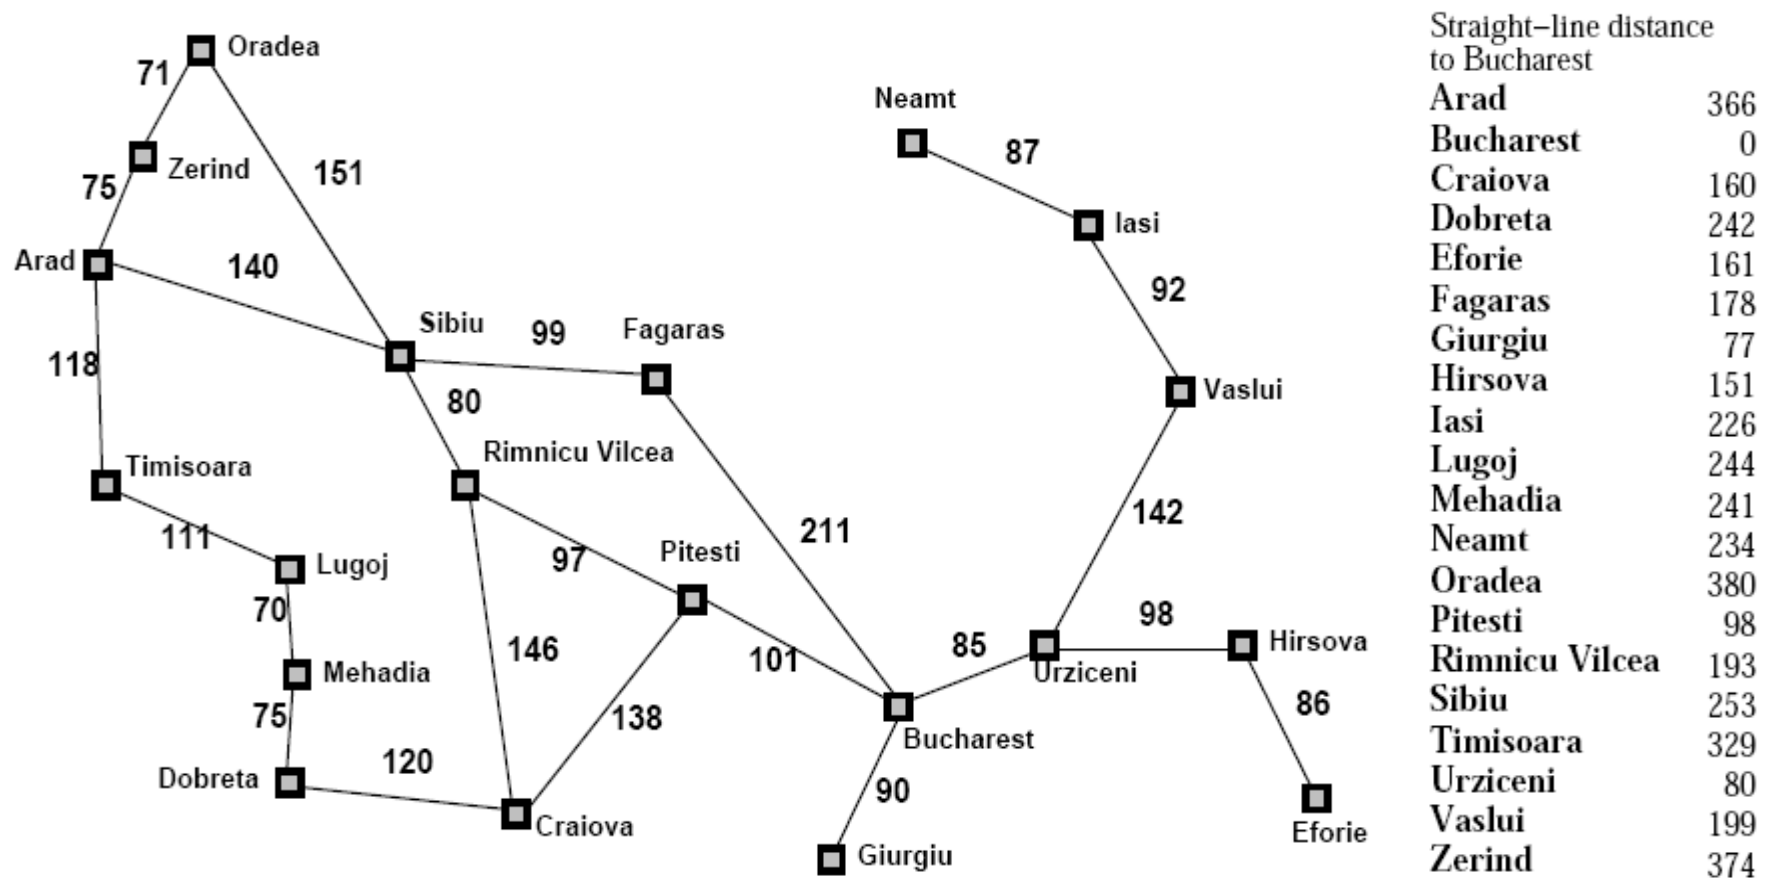
9

- Strategy: expand nodes which appear closest to goal
  - ▣ Heuristic (启发函数) : function which maps states to distance
- A common case:
  - ▣ Best-first takes you straight to the (wrong) goal
- Worst-case: like a badly-guided DFS



# Romania with Step Costs in Km

10



# Greedy Search

11

Evaluation function  $f(n) = h(n)$  (heuristic function 启发函数)

= estimate of cost from  $n$  to the closest goal

(节点 $n$ 到目标节点的最低耗散路径的耗散估计值)

- ▣  $h(n)$  a problem-specific function
- ▣  $h(n) = 0$  if  $n$  is goal node

E.g.,  $h_{\text{SLD}}(n)$  = straight-line distance from  $n$  to Bucharest

Greedy search expands the node that **appears** to be closest to goal (试图扩展离目标节点最近的点)


# Greedy Search Example

12



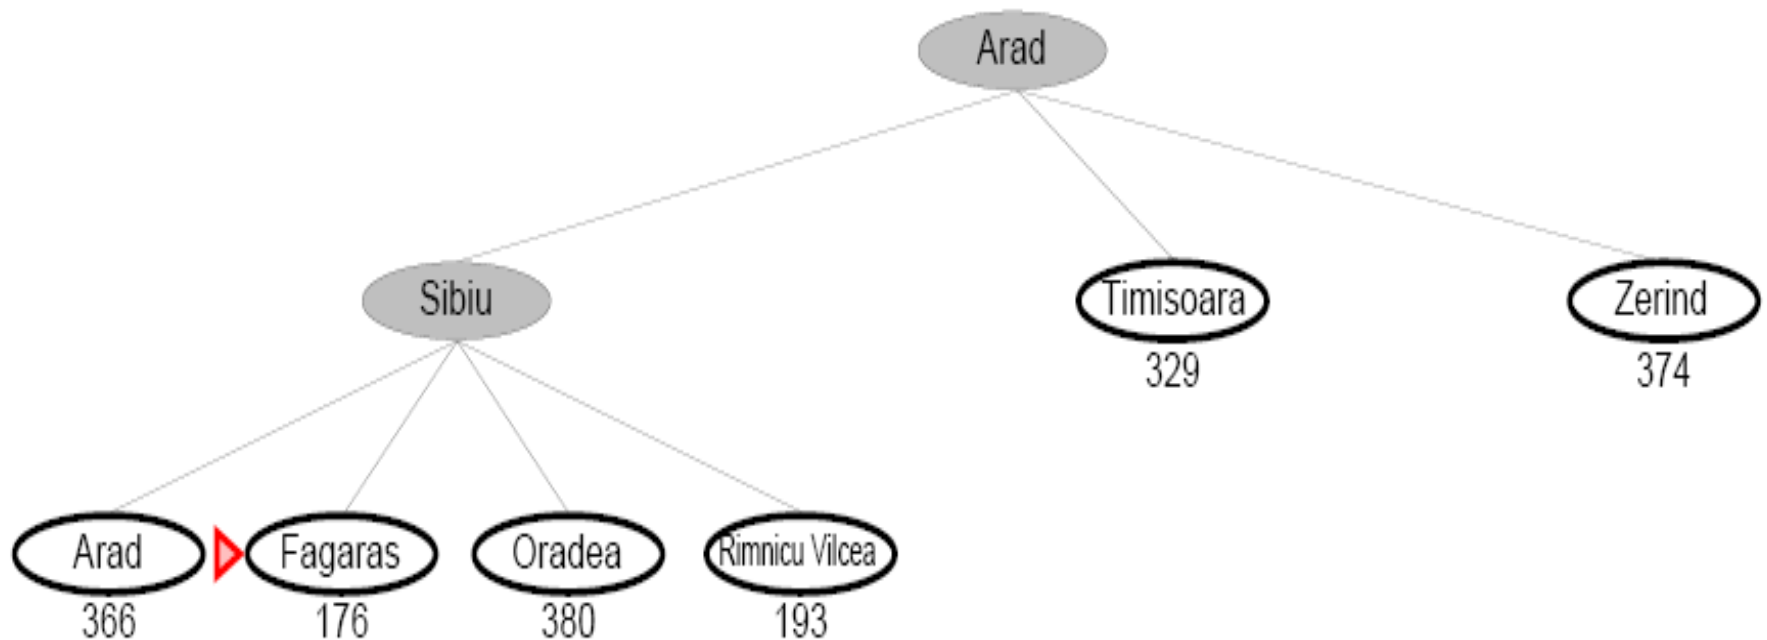
# Greedy Search Example

13



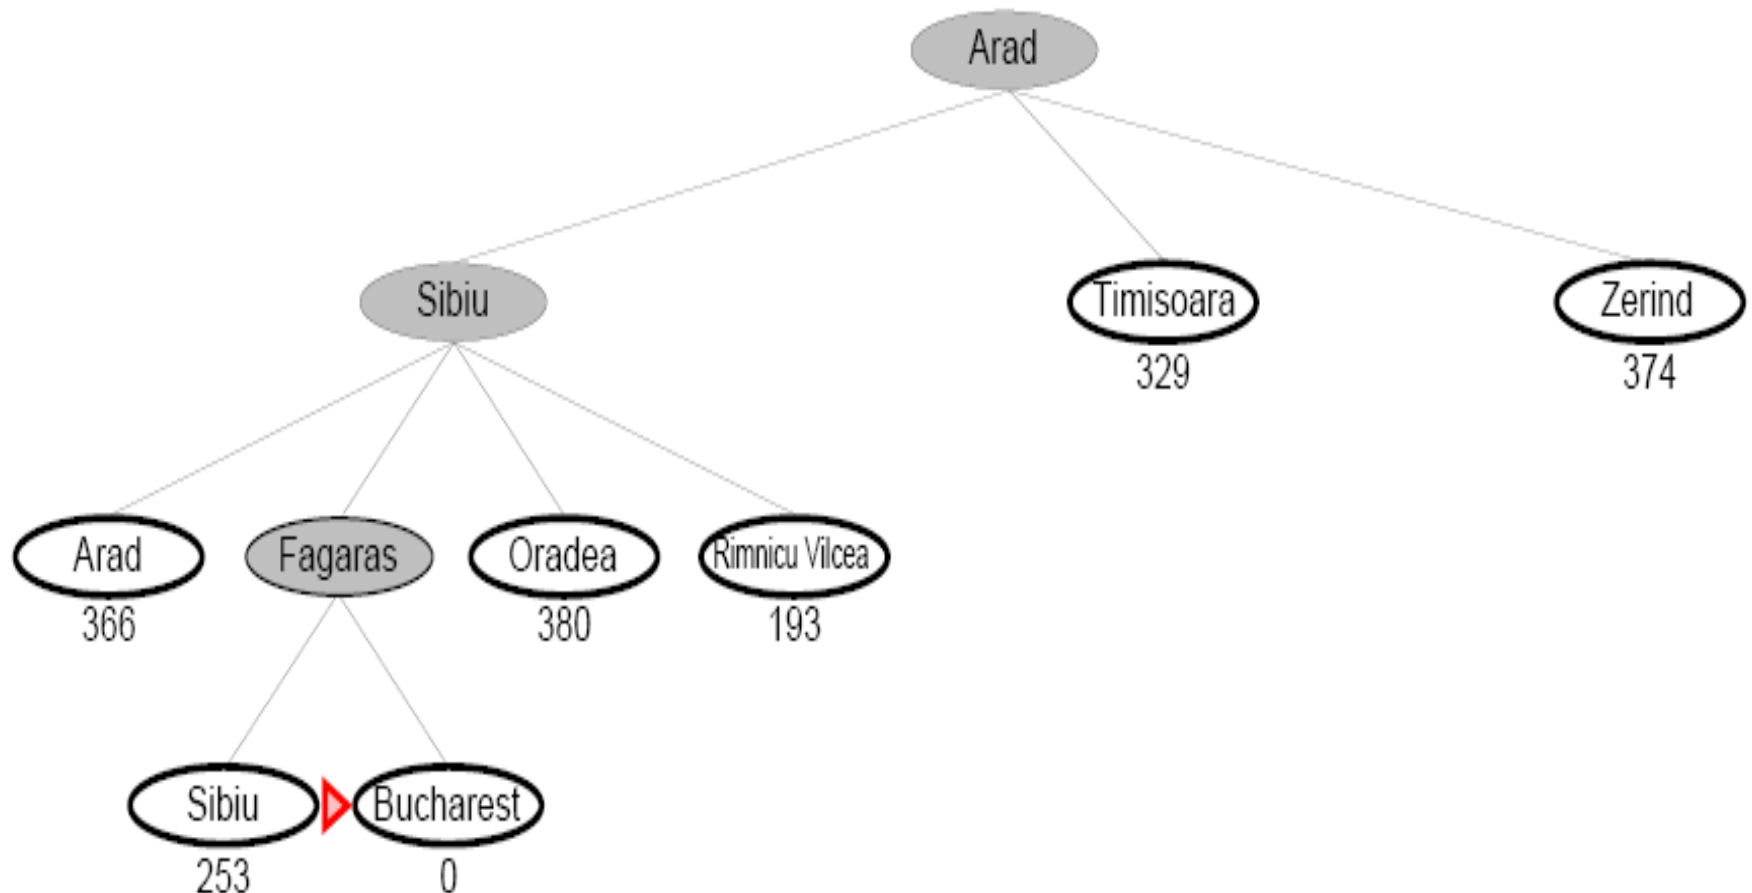
# Greedy Search Example

14



# Greedy Search Example

15



# Properties of Greedy Search

16

Complete??



# Properties of Greedy Search

17

Complete?? No — can get stuck in loops, e.g., with Oradea as goal,

Iasi → Neamt → Iasi → Neamt →

Complete in finite space with repeated-state checking

Time??

# Properties of Greedy Search

18

Complete?? No — can get stuck in loops, e.g., with Oradea as goal,

Iasi → Neamt → Iasi → Neamt →

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??

b: Branching factor

d: Solution depth

m: Maximum depth

# Properties of Greedy Search

19

Complete?? No — can get stuck in loops, e.g., with Oradea as goal,

Iasi → Neamt → Iasi → Neamt →

Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??  $O(b^m)$  — keeps all nodes in memory

b: Branching factor  
d: Solution depth  
m: Maximum depth

Optimal??

# Properties of Greedy Search

20

Complete?? No — can get stuck in loops, e.g., with Oradea as goal,

Iasi  $\rightarrow$  Neamt  $\rightarrow$  Iasi  $\rightarrow$  Neamt  $\rightarrow$


Complete in finite space with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??  $O(b^m)$  — keeps all nodes in memory

b: Branching factor  
d: Solution depth  
m: Maximum depth


Optimal?? No



# Combining UCS and Greedy

21

- **Uniform-cost** orders by path cost, or *backward cost*  $g(n)$
- **Greedy-search** orders by goal proximity, or *forward cost*  $h(n)$



- **A\* Search** orders by the sum:  $f(n) = g(n) + h(n)$

# A\* Search

22

Evaluation function  $f(n) = g(n) + h(n)$

$g(n)$  = cost so far to reach  $n$  — 到达节点  $n$  的耗散

$h(n)$  = estimated cost to goal from  $n$

— 启发函数：从节点  $n$  到目标节点的最低耗散路径的耗散估计值

$f(n)$  = estimated total cost of path through  $n$  to goal

— 经过节点  $n$  的最低耗散的估计函数

A\* search uses an **admissible** heuristic 可采纳启发式

i.e.,  $h(n) \leq h^*(n)$  where  $h^*(n)$  is the **true** cost from  $n$ .

(Also require  $h(n) \geq 0$ , so  $h(G) = 0$  for any goal  $G$ .)



**Theorem:** A\* search is optimal

E.g.,  $h_{\text{SLD}}(n)$  never overestimates the actual road distance (**SLD: Straight-Line Distance**)


# A\* Search Example

23

▶ Arad  
 $366 = 0 + 366$

# A\* Search Example

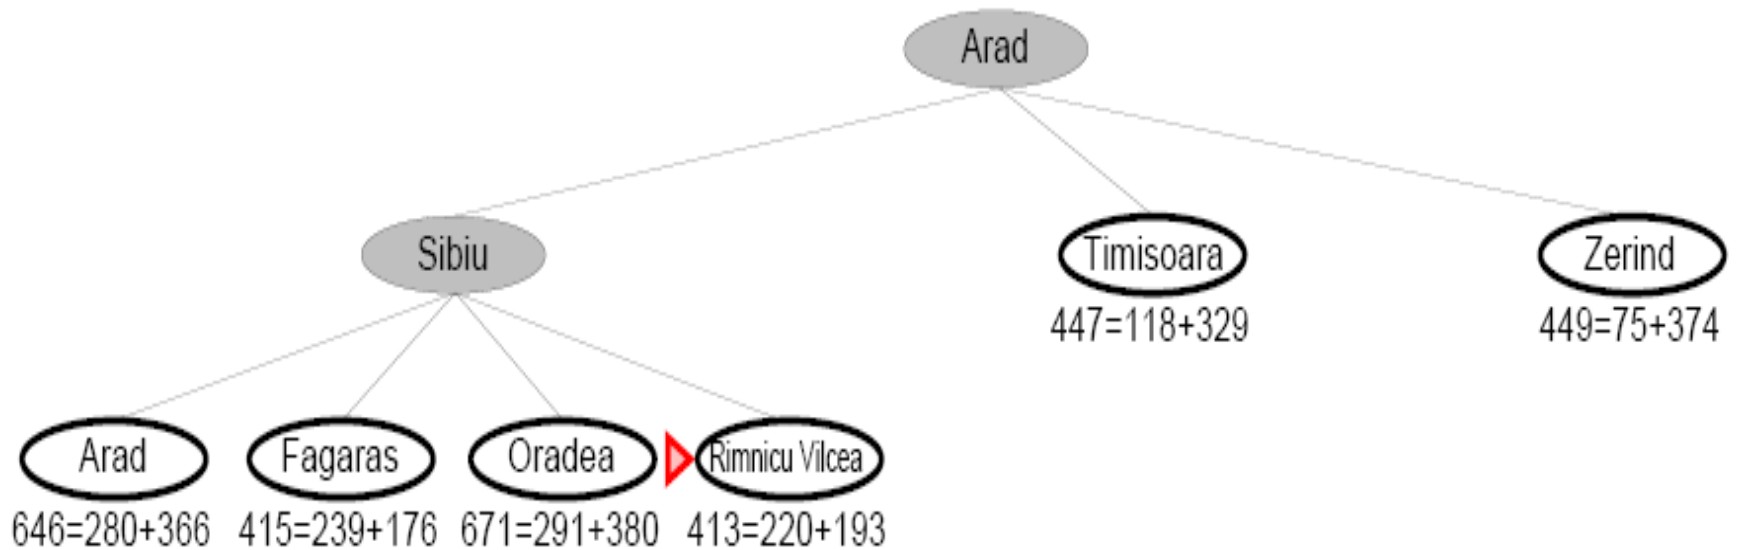
24





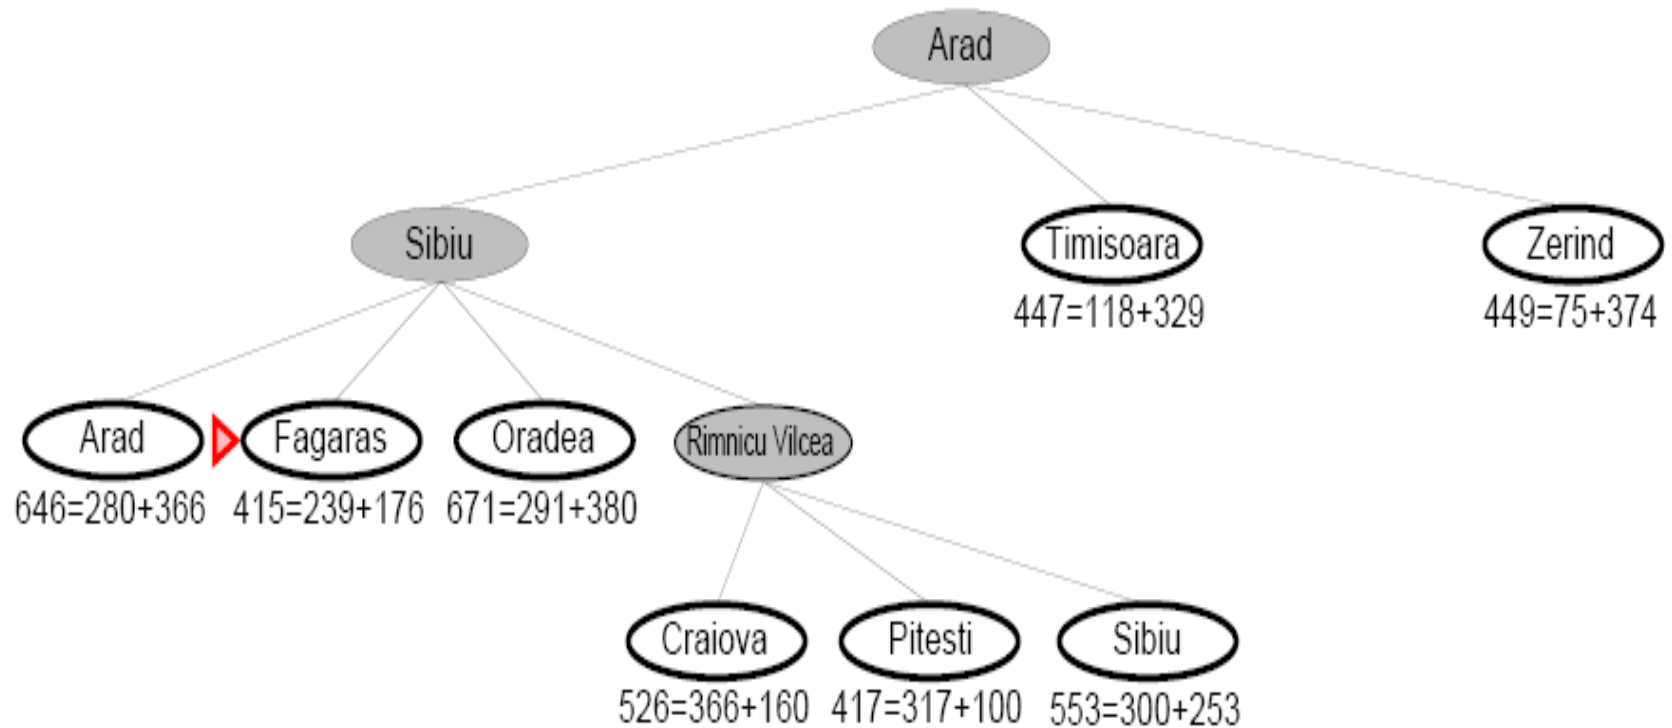
# A\* Search Example

25



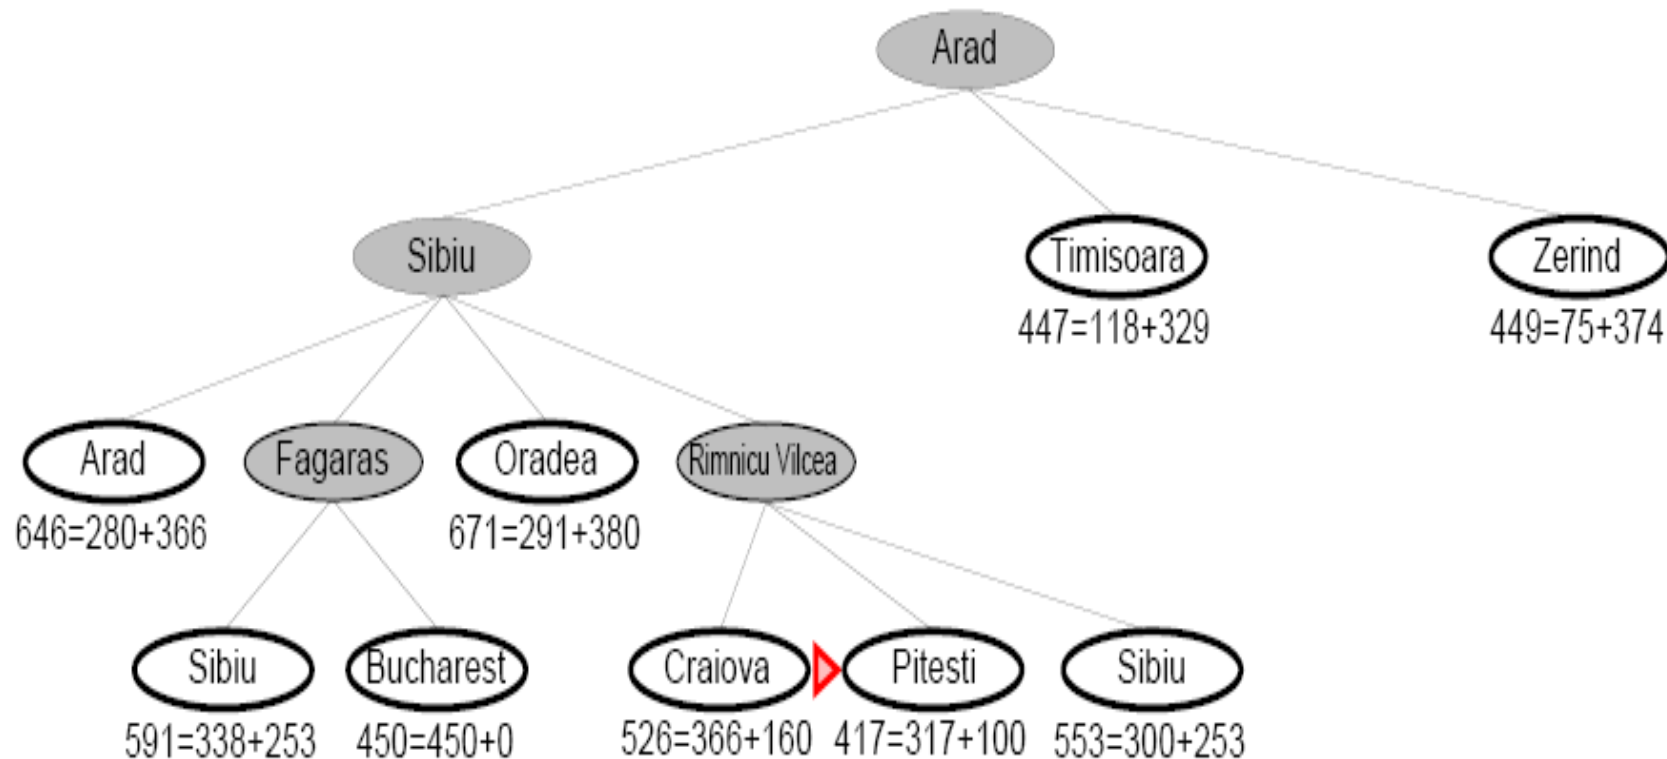
# A\* Search Example

26



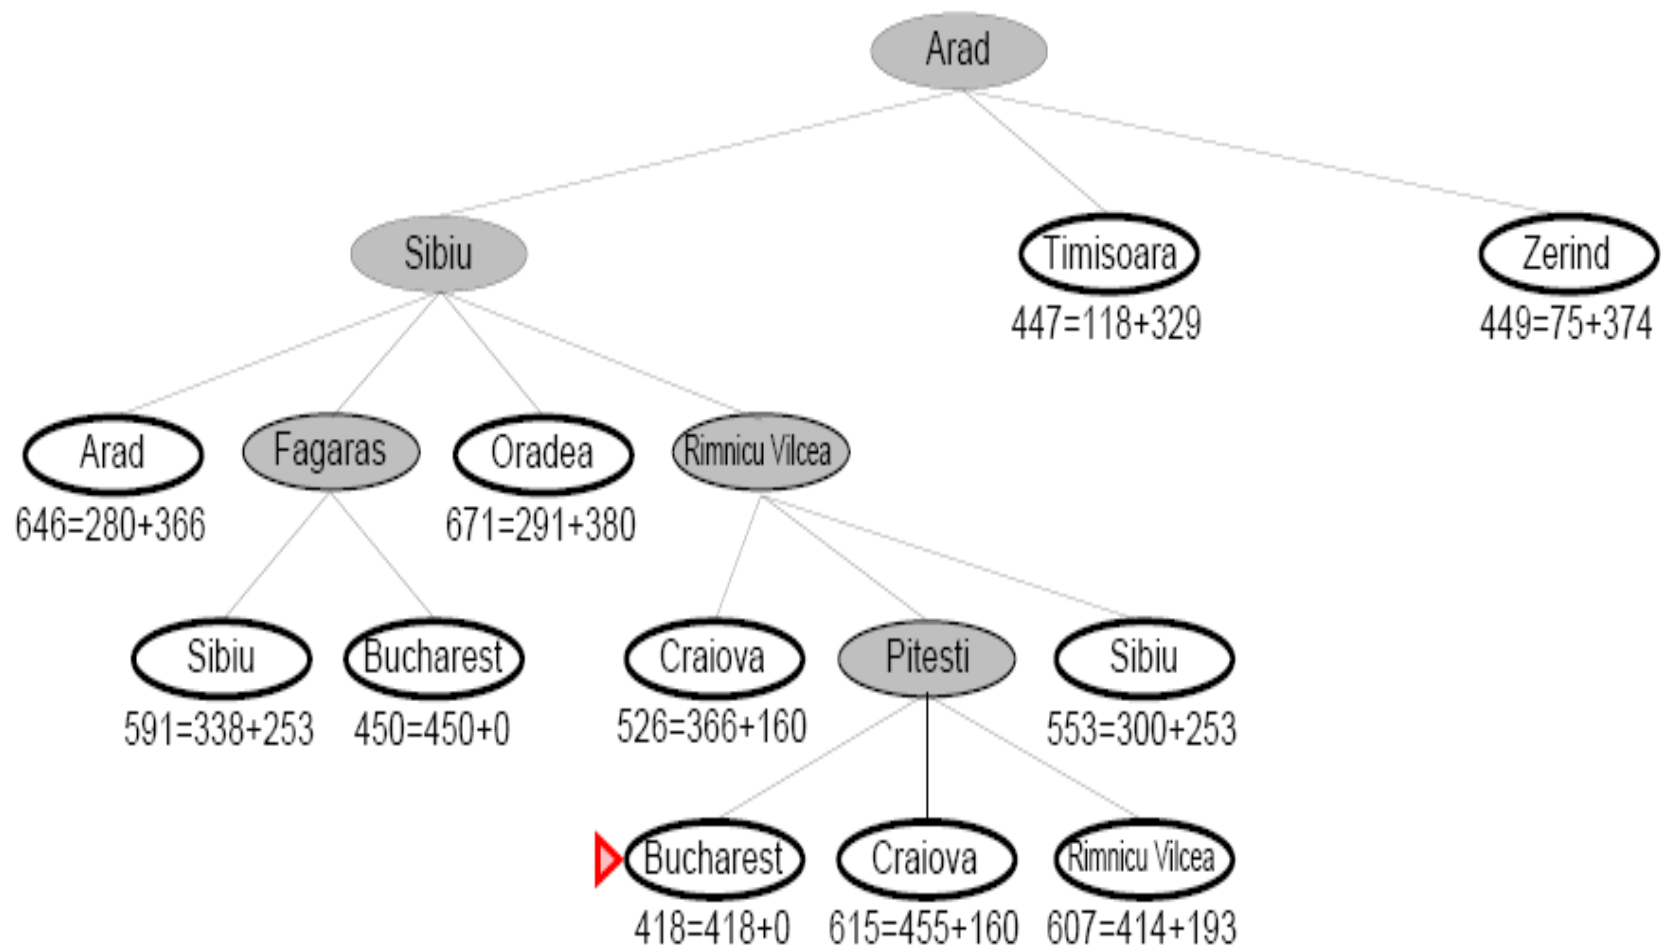
# A\* Search Example

27



# A\* Search Example

28



# Admissible Heuristics

29


- A\* heuristic  $h(n)$  is **admissible** (可采纳的) if for every node  $n$ ,  $h(n) \leq h^*(n)$ , where  $h^*(n)$  is the **true** cost to reach the goal state from  $n$ .
- An admissible heuristic **never overestimates** the cost to reach the goal (从不会过高估计到达目标的耗散), i.e., it is **optimistic** (乐观的)
- Example:  $h_{SLD}(n)$  (never overestimates the actual road distance)
- **Theorem**: If  $h(n)$  is admissible, A\* using TREE-SEARCH is optimal

Coming up with admissible heuristics is most of what's involved in using A\* in practice

# Optimality of $A^*$ (proof)

30

Suppose some suboptimal (非最优) goal  $G_2$  has been generated and is in the fringe. Let  $n$  be an unexpanded node in the fringe such that  $n$  is on a shortest path to an optimal goal  $G$ .




- $f(G_2) = g(G_2)$                       since  $h(G_2) = 0$
- $g(G_2) > g(G)$                       since  $G_2$  is suboptimal
- $f(G) = g(G)$                       since  $h(G) = 0$
- $f(G_2) > f(G)$                       from above

# Optimality of $A^*$ (proof)

31

Suppose some suboptimal goal  $G_2$  has been generated and is in the fringe. Let  $n$  be an unexpanded node in the fringe such that  $n$  is on a shortest path to an optimal goal  $G$ .



- $f(G_2) > f(G)$  from above
- $h(n) \leq h^*(n)$  since  $h$  is admissible
- $g(n) + h(n) \leq g(n) + h^*(n) = g(G) = f(G)$   $n$  is on a shortest path to an optimal goal  $G$
- $f(n) \leq f(G)$

Hence  $f(G_2) > f(n)$ , and  $A^*$  will never select  $G_2$  for expansion

# Consistent Heuristics

32


- A\* heuristic is **consistent** (一致) if for every node  $n$ , every successor  $n'$  of  $n$  generated by any action  $a$ ,

$$h(n) \leq c(n,a,n') + h(n')$$

- If  $h$  is consistent, we have

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$

- i.e.,  $f(n)$  is non-decreasing along any path.
- **Theorem:** If  $h(n)$  is consistent, A\* using GRAPH-SEARCH is optimal






# Optimality of $A^*$

33

- $A^*$  expands nodes in order of increasing  $f$  value
- Gradually adds " $f$ -contours (等值线)" of nodes
- Contour  $i$  has all nodes with  $f=f_i$ , where  $f_i < f_{i+1}$



# Properties of A\*

34

- Complete? Yes (unless there are infinitely many nodes with  $f \leq f(G)$  )
- Time? A\*算法对于任何给定的启发函数都是效率最优的  
But still exponential
- Space? Keeps all nodes in memory
- Optimal? Yes

A\* expands all nodes with  $f(n) < C^*$

A\* expands some nodes with  $f(n) = C^*$

A\* expands no nodes with  $f(n) > C^*$

# 8-puzzle Revisited

35

- **8-puzzle** — 把棋子水平或者竖直地滑动到空格中，直到目标局面：

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- 平均解步数是**22**步。分支因子约为3
  - ▣ 到达深度为22的穷举搜索将考虑约 $3^{22} \approx 3.1 \times 10^{10}$
  - ▣ 状态个数 $O((n+1)!)$ ，NP完全问题

# Admissible Heuristics

36

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles (错位的棋子数)
- $h_2(n)$  = total Manhattan distance (所有棋子到其目标位置的水平和垂直距离和)  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

# Admissible Heuristics

37

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles (错位的棋子数)
- $h_2(n)$  = total Manhattan distance (所有棋子到其目标位置的水平和垂直距离和)  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$  8
- $h_2(S) = ?$   $3+1+2+2+2+3+3+2 = 18$

# Dominance

38

- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible)
- then  $h_2$  **dominates**  $h_1$  (dominate 统治、占优)
- $h_2$  is better for search
- Typical search costs (average number of nodes expanded):
  - $d=12$       IDS = 3,644,035 nodes  
                   $A^*(h_1) = 227$  nodes  
                   $A^*(h_2) = 73$  nodes
  - $d=24$       IDS = too many nodes  
                   $A^*(h_1) = 39,135$  nodes  
                   $A^*(h_2) = 1,641$  nodes


Given any admissible heuristics  $h_a, h_b$ ,

$$h(n) = \max(h_a(n); h_b(n))$$

is also admissible and dominates  $h_a, h_b$

With  $A^*$ : a trade-off between quality of estimate and work per node!

使用  $f(n) = g(n) + h(n)$  搜索  
 $h(n)$  = 错放数字个数  
 $g(n)$  = 节点深度



Most of the work is in coming up with admissible heuristics

How do we get good heuristics?

Just relax...







# Relaxation

41

Goal: move from triangle to circle



Hard



Easy

Heuristic:

$$h(s) = \text{ManhattanDistance}(s, (2, 5))$$

□ Interpretation: relax → removing constraints

# Relaxed Problems

42

- A problem with fewer restrictions on the actions is called a **relaxed problem** (松弛问题)
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem  
一个松弛问题的最优解的耗散是原问题的一个可采纳的启发式
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  gives the shortest solution  
如果棋子可以任意移动，则  $h_1$  给出最短的确切步数
- If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  gives the shortest solution  
如果棋子可以移动到任意相邻的位置，则  $h_2$  给出最短的确切步数

**Key point:** the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem

# Relaxed Problems

43

## □ 构造松弛问题

- 原问题：一个棋子可以从方格A移动到方格B，如果A与B水平或者垂直相邻**而且**B是空的
- 松弛1：一个棋子可以从方格A移动到方格B，如果A与B相邻 —  $h_2$
- 松弛2：一个棋子可以从方格A移动到方格B，如果B是空的
- 松弛3：一个棋子可以从方格A移动到方格B —  $h_1$

## □ 如果有一个可采纳启发式的集合 $\{h_1, \dots, h_m\}$

$$h(n) = \max(h_1(n), \dots, h_m(n))$$

可采纳并比成员启发式更有优势

# Evaluation Function $f(n)$

44

$h(n)$  — heuristic, estimate of cost from  $n$  to the closest goal  
(节点  $n$  到目标节点的最低耗散路径的耗散估计值)

$g(n)$  — path cost to  $n$  (初始节点到这个节点的路径损耗的总和)

Possible evaluation functions:

- $f(n) = g(n)$   $\equiv$  **Uniform Cost**
- $f(n) = h(n)$   $\equiv$  **Greedy**
- $f(n) = g(n) + h(n)$   $\equiv$  **A\***  
*estimates the total cost of a solution path  
which goes through node  $n$*

# Informed (Heuristic) Search

45

- Summary of uninformed tree (problem-solving) search:
  - ▣ Algorithms differ by method of *expansion*, or *choice of state-action* sequence for offline evaluation
  - ▣ Complexity tradeoffs, but poor (worst case or typical case) performance from all when state space is large
- Improvement:
  - ▣ Exploit knowledge about which successor states are preferable
  - ▣ “Best-First” search: expand nodes based on evaluation function  $f(n)$
  - ▣  $f(n) \sim$  estimate of distance to goal
  - ▣ expand node  $n$  with lowest evaluation
  - ▣ “Best-First” algorithms characterized by choice of evaluation function

# Outline


46

- Best-first search (最佳优先搜索)
  - ▣ Greedy best-first search
  - ▣ A\* search
- Heuristics
- Local search algorithms (局部搜索算法)
  - ▣ Hill-climbing search (爬山法搜索)
  - ▣ Simulated annealing search (模拟退火搜索)
  - ▣ Local beam search (局部剪枝搜索)
  - ▣ Genetic algorithms (遗传算法)

# Large Scale Problems with A\*

47

- What states get expanded?
  - ▣ All states with f-cost less than optimal goal cost
- In huge problems, often A\* isn't enough
  - ▣ State space just too big
  - ▣ Can't visit all states with f less than optimal
  - ▣ Often, can't even store the entire fringe
- Solutions
  - ▣ Better heuristics
  - ▣ Greedy hill-climbing (fringe size = 1)
  - ▣ Beam search (limited fringe size)



# Limited Memory Options

48

- Bottleneck: not enough memory to store entire fringe
- Hill-Climbing Search:
  - ▣ Only “best” node kept around, no fringe!
  - ▣ Usually prioritize successor choice by  $h$  (greedy hill climbing)
  - ▣ Compare to greedy backtracking, which still has fringe
- Beam Search (Limited Memory Search)
  - ▣ In between: keep  $K$  nodes in fringe
  - ▣ Dump lowest priority nodes as needed
  - ▣ Can prioritize by  $h$  alone (greedy beam search), or  $h+g$  (limited memory  $A^*$ )
- No guarantees once you limit the fringe size!



# Local Search Algorithms


49

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of “complete” configurations (完全状态)
  - ▣ Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- keep a single “current” state, try to improve it
  - ▣ until you can’t make it better
- Constant space, suitable for online as well as offline search
- Generally much more efficient (but incomplete)



# Example: $n$ -Queens

50



- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal



$h = 5$



$h = 2$



$h = 0$

# Optimization Problems

51

- Now a different setting:
  - ▣ Each state  $s$  has a score or cost,  $f(s)$ , that we can compute
  - ▣ The goal is to find the state with the highest (or lowest) score, or a reasonably high (low) score
  - ▣ We do not care about the path
  - ▣ This is an optimization problem
  - ▣ Enumerating the states is intractable
  - ▣ Previous search algorithms are too expensive
  - ▣ No known algorithm for finding optimal solution efficiently

# Hill-Climbing Search

52

- Simple, general idea:
  - ▣ Start wherever
  - ▣ Always choose the best neighbor
  - ▣ If no neighbors have better scores than current, quit
- Why can this be a terrible idea?
  - ▣ Complete?
  - ▣ Optimal?
- What's good about it?

# Hill-Climbing Search

53

- “Like climbing Everest in thick fog with amnesia (健忘症)”
  - ▣ Can’t see the entire landscape all at once


```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```

# Hill-Climbing Search

54

- Problem: depending on initial state, can get stuck in local maxima (局部最大值)



Random-restart hill climbing overcomes local maxima — trivially complete  
Random sideways moves escape from shoulders 😊 loop on at maxima ☹️

# Hill-Climbing with Random Restarts

55

## □ Very simple modification:

1. When stuck, pick a random new starting state and re-run hill-climbing from there
2. Repeat this  $k$  times
3. Return the best of the  $k$  local optima

- ▣ Can be very effective
- ▣ Should be tried whenever hill-climbing is used
- ▣ Fast, easy to implement; works well for many applications where the solution space surface is not too “bumpy” (i.e., not too many local maxima)

# Hill-Climbing Search: 8-Queens Problem

56


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18

- $h$  = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$  for the above state



# Hill-Climbing Search: 8-Queens Problem

57



- A local minimum with  $h = 1$

# Life Lesson

58

**Sometimes one needs to temporarily step backward in order to move forward**

- Lesson applied to iterative, local search:
  - Sometimes one needs to move to an *inferior neighbor* in order to escape a local optimum

# Simulated Annealing Search

## (模拟退火搜索)

59

- Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to "temperature"
  local variables: current, a node
                   next, a node
                   T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to  $\infty$  do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E$  ← VALUE[next] - VALUE[current]
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

# Properties of Simulated Annealing Search

60

- One can prove: If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Widely used in VLSI layout, airline scheduling, etc

# Local Beam Search (局部剪枝搜索)

61

Keep track of  $k$  states rather than just one


Start with  $k$  randomly generated states

At each iteration, all the successors of all  $k$  states are generated


If any one is a goal state, stop; else select the  $k$  best successors from the complete list and repeat.

# Local Beam Search

- Like greedy search, but keep  $K$  states at all times:



Greedy Search



Beam Search

- The best choice in MANY practical settings
- Not the same as  $k$  searches run in parallel!
- Searches that find good states recruit other searches to join them
- Variables: beam size, encourage diversity?
- **Problem:** quite often, all  $k$  states end up on same local hill
- **Idea:** choose  $k$  successors randomly, biased towards good ones

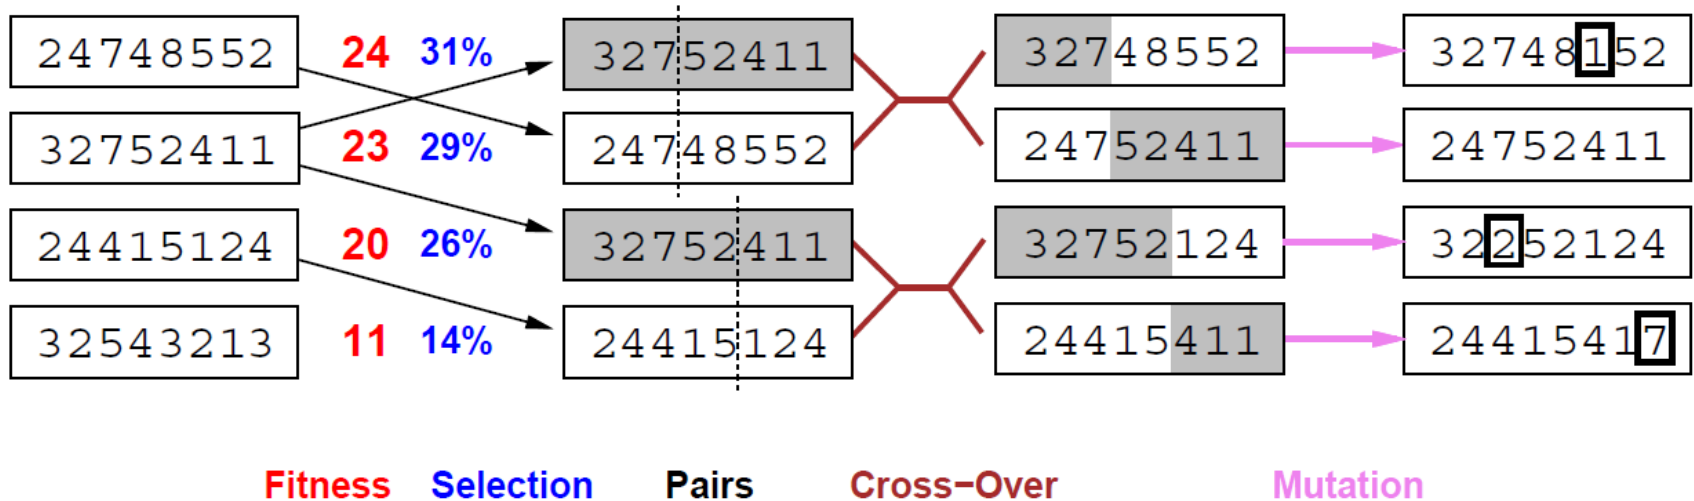
# Genetic Algorithms

63

- Genetic algorithms use a natural selection metaphor
- A successor state is generated by combining two parent states
- Start with  $k$  randomly generated states (**population**种群)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (**fitness function**适应度函数). Higher values for better states.
- Produce the next generation of states by selection, crossover, and mutation (选择, 杂交, 变异)

# Genetic Algorithms

64



- Fitness function: number of non-attacking pairs of queens 不互相攻击的皇后对的数目 (min = 0, max =  $8 \times 7/2 = 28$ )


$$24/(24+23+20+11) = 31\%$$

$$23/(24+23+20+11) = 29\% \text{ etc}$$



# Genetic Algorithms

65



# Summary 1

66

Heuristic functions estimate costs of shortest paths

Good heuristics can dramatically reduce search cost

Greedy best-first search expands lowest  $h$

- incomplete and not always optimal

A\* search expands lowest  $g + h$

- complete and optimal

- also optimally efficient (up to tie-breaks, for forward search)

Admissible heuristics can be derived from exact solution of relaxed problems

# Summary2

67

## □ Local search algorithms

the **path** to the goal is irrelevant; the goal state itself is the solution

keep a single "current" state, try to improve it

### ▣ Hill-climbing search

- depending on initial state, can get stuck in local maxima

### ▣ Simulated annealing search

- escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency

### ▣ Local beam search

- Keep track of  $k$  states rather than just one

### ▣ Genetic algorithms

# Uninformed/Informed Search

## — Main points

68

- 好的启发式搜索能大大提高搜索性能
- 但由于启发式搜索需要抽取与问题本身有关的特征信息，而这种特征信息的抽取有时会比较困难，因此盲目搜索仍不失为一种有用的搜索策略。

# Uninformed/Informed Search

## — Main points

69

- 好的搜索策略应该
  - ▣ 引起运动—避免原地踏步
  - ▣ 系统—避免兜圈
  - ▣ 运用启发函数—缓解组合爆炸

# Uninformed/Informed Search

## — Main points

70

- 搜索树 vs 搜索图
  - ▣ 搜索树：结点有重复，但登记过程简单
  - ▣ 搜索图：结点无重复，但登记过程复杂（每次都要查重）
    - 省空间，费时间。

# 作业

71

□ 4.1, 4.2, 4.6, 4.7

▣ 习题编号按人民邮电出版社第二版

